

Stop Chasing AI. Start Chasing Business Impact.

AI has moved from experimentation to boardroom priority.

Generative AI. Copilots. AI agents. Agentic AI. Autonomous workflows.

Every few months, there is a new technology that enterprises are told they cannot afford to ignore.

But there is a more important question that often gets lost:

What business problem are we actually solving?

Because the companies creating the most value from AI are not necessarily the ones deploying the most sophisticated technology.

They are the ones getting better at **matching the technology to the problem.**

The AI adoption paradox

AI adoption is everywhere. AI impact is not.

McKinsey's 2025 State of AI research found that **88% of organizations are regularly using AI in at least one business function**, yet only **39% report any enterprise-level EBIT impact** from AI. Nearly two-thirds have not yet started scaling AI across the organization.

BCG's research tells a similar story: only around **5% of companies are currently "future-built" for AI**, while roughly 60% report little material value despite significant investment.

The problem isn't adoption.

The problem is value creation.

The companies getting it right start with the business problem

Consider how some leading companies are approaching AI.

Amazon: AI for the supply chain, not AI for AI's sake

Amazon is using AI to solve one of retail's hardest problems: **putting the right inventory in the right place at the right time.**

Its AI-powered forecasting models predict demand across hundreds of millions of products.

The reported results include:

- **20% improvement in regional forecasting accuracy** for millions of popular items
- **10% improvement in long-term national forecasts** for major events

- Better inventory placement
- Faster delivery
- Fewer unnecessary transportation miles

The lesson?

Amazon didn't start with:

“Where can we deploy an AI agent?”

It started with:

“**How do we make our supply chain more intelligent?**”

The technology followed the problem.

Walmart: AI where optimization matters

Walmart uses AI across logistics to optimize routes, trailer loading and inventory movement.

Its Route Optimization technology has reportedly:

- Eliminated **30 million unnecessary miles**
- Avoided **94 million pounds of CO₂**
- Optimized more than **110,000 inefficient routes**

More recently, Walmart has expanded AI into pallet building, load planning, predictive maintenance and fulfillment orchestration.

Notice something important:

There is no need to make every interaction conversational.

Sometimes the best AI experience is **invisible**.

A truck gets a better route.

A trailer gets loaded better.

A machine gets serviced before it fails.

The employee may never even know that AI was involved.

And that's perfectly fine.

JPMorgan: Measure productivity, not model sophistication

JPMorgan Chase has spent years applying AI and machine learning across areas such as fraud, credit, personalization and software development.

Its 2025 reporting highlights measurable productivity gains and AI-powered tools embedded directly into employee workflows. The bank says its AI capabilities are producing tangible financial results, while its engineering organization has reported significant efficiency improvements from AI coding tools.

The important principle is simple:

AI isn't the product. Productivity is.

If an engineer can complete work faster, that is the value.

If an advisor can spend more time with a client, that is the value.

If a risk team can identify problems earlier, that is the value.

The model behind it is secondary.

Morgan Stanley: Put AI inside the workflow

Morgan Stanley provides another useful example.

Rather than creating another destination that employees have to visit, the firm has embedded AI into existing workflows.

Its research teams use generative AI to search and synthesize tens of thousands of proprietary research reports. Its wealth-management teams use AI to summarize meetings, surface research and prepare for client interactions.

The principle is powerful:

Don't ask employees to "use AI."

Redesign the workflow so that **AI naturally becomes part of how the work gets done.**

The technology should follow the problem

This is where enterprises need to become much more disciplined.

Not every problem needs an LLM.

Not every LLM problem needs an agent.

Not every agent needs autonomy.

A simple way to think about it:

Business problem	Right technology may be
Demand forecasting	Statistical models / ML
Fraud detection	ML / anomaly detection
Equipment failure	Predictive AI
Document search	RAG + LLM
Content generation	Generative AI
Customer support	GenAI + workflow automation
Complex multi-step process	AI agents
Dynamic operational decisions	Agentic AI
Simple repetitive process	Traditional automation

The question should never be:

“How do we use the latest AI?”

It should be:

“What is the simplest technology that can create a meaningful improvement in this business process?”

Sometimes that will be an advanced AI agent.

Sometimes it will be a machine-learning model.

Sometimes it will be a rules engine.

And sometimes the right answer will be:

Don't use AI.

From AI projects to AI value pools

The companies that are getting this right are changing how they prioritize AI.

Instead of starting with technology, they start with **value pools**.

1. Find the economic problem

Where are you losing money?

Where are customers leaving?

Where are employees spending too much time?

Where is working capital trapped?

Where are decisions too slow?

2. Quantify the opportunity

Put a number against it.

- ₹10 Cr inventory inefficiency
- 50,000 hours of manual work
- 8% customer churn
- 15% forecast error
- 20% machine downtime

Without a baseline, there is no meaningful AI ROI.

3. Choose the appropriate technology

Only now decide whether you need:

Automation → ML → GenAI → Agents → Agentic AI

Technology should be a consequence of the business requirement.

4. Redesign the workflow

This is where many AI initiatives fail.

McKinsey's research found that **workflow redesign is one of the strongest factors associated with achieving bottom-line impact from generative AI.**

Putting an AI tool on top of a broken process doesn't transform the process.

It simply makes the broken process faster.

5. Measure the outcome

Every AI initiative should have a business metric attached to it.

Not:

- Number of prompts
- Number of users
- Number of AI agents
- Number of models deployed

But:

- Revenue generated
- Cost eliminated
- Hours saved
- Forecast accuracy improved
- Downtime reduced
- Conversion increased
- Risk reduced
- Customer satisfaction improved

If you cannot connect the AI initiative to a business metric, you probably don't have an AI strategy yet.

The next phase of enterprise AI

The first phase of enterprise AI was about **experimentation**.

The second phase was about **adoption**.

The next phase will be about **economic value**.

That will require CIOs and business leaders to become much more selective.

Not:

“Where can we put AI?”

But:

“Where can AI materially change the economics of our business?”

That shift will separate AI leaders from AI followers.

The winners won't necessarily have the biggest models.

They won't necessarily have the most agents.

They won't necessarily have the largest AI budgets.

They will know where AI belongs—and, equally importantly, where it doesn't.

AI strategy is not about using more AI.

It is about using the right AI, in the right workflow, against the right problem, with a measurable business outcome.

That is where the hype ends.

And where real AI transformation begins.